

# Assignment 1 - Yang Chen

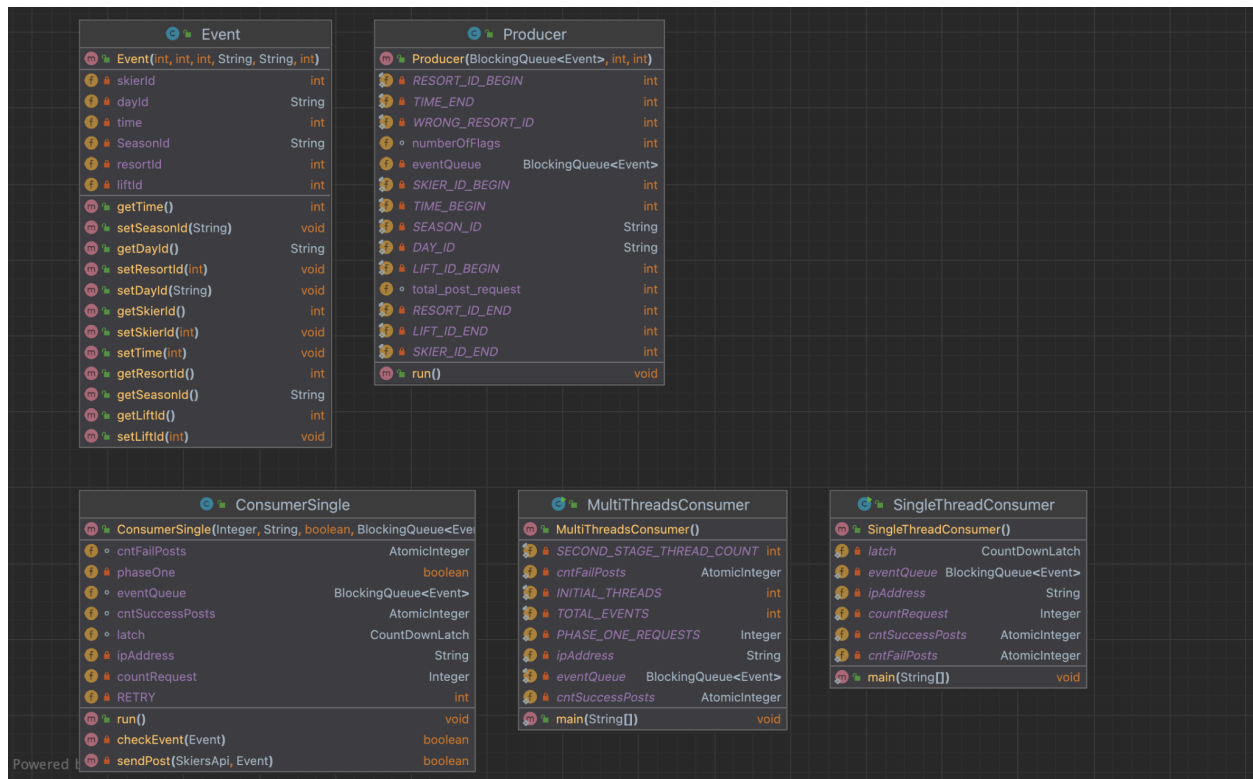
## Git repo

<https://github.com/Yangyanggogo/SkiResortDistributedSystem>

## Client design description

### Client Part 1

Here is my client-part 1 UML diagram, which presents the classes included in client-part 1.



I implemented the client-part 1 using a general design pattern of producer-consumer. The **producer** is a single thread, responsible for generating the events and storing them in a **BlockingQueue** (thread-safe), which consumers can access.

The **ConsumerSingle** class is responsible for consuming the event (defined in **Event class**) in the BlockingQueue and sending a post request to my server, with an event in each request's body.

According to the assignment's requirements, phase one has 32 threads and each thread sends 1000 post requests to the server. As soon as one of these 32 threads finishes its task, I can allocate x number of threads and start phase two (x is a variable to test my multithreads client-part 1).

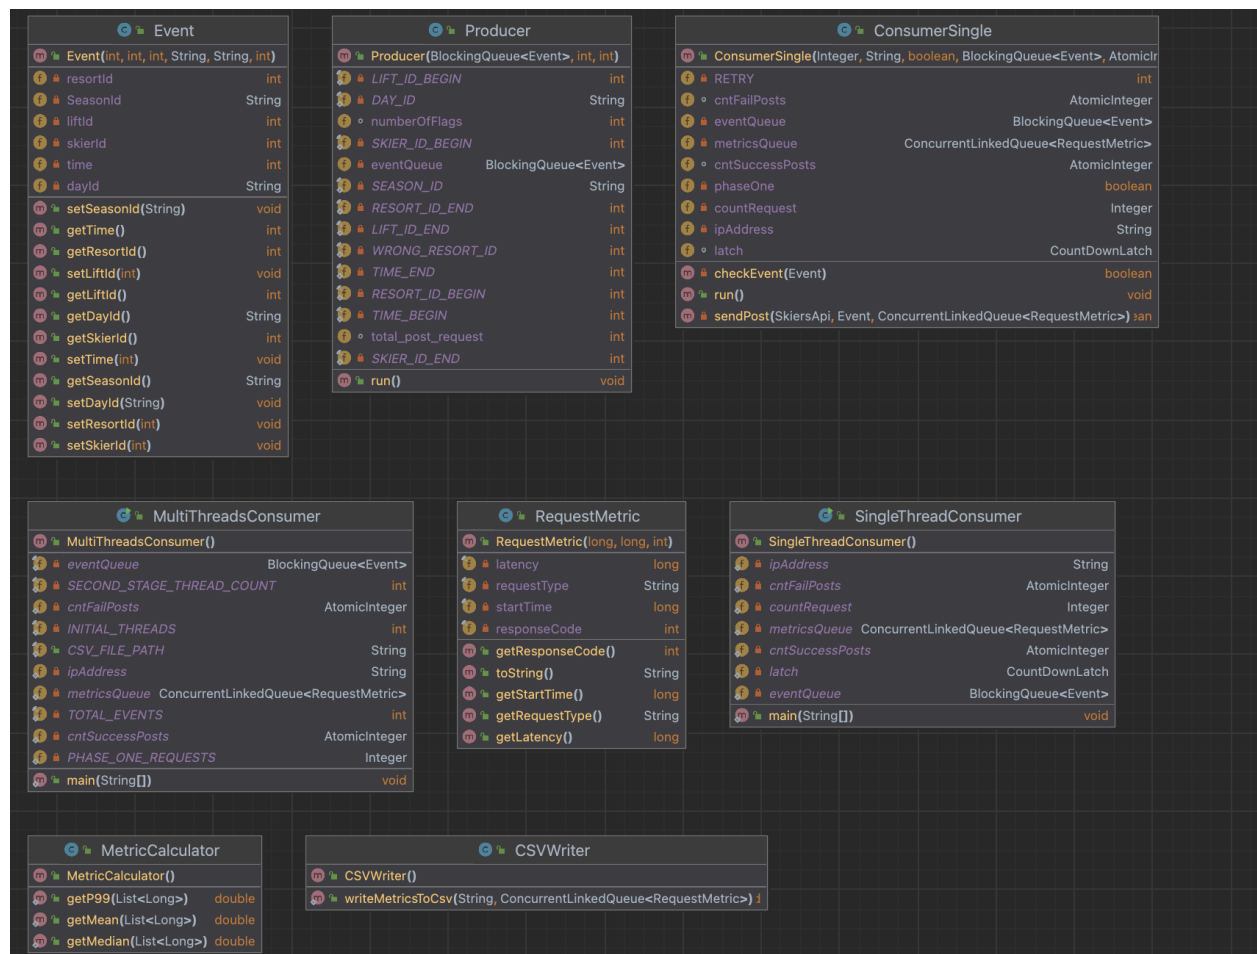
My ConsumerSingle class can handle both requirements by passing the boolean "phaseOne" to this class.

32 threads in phase one can be terminated when finishing sending 1000 requests. To terminate my threads in phase two, I added extra "invalid events" in the BlockingQueue, when any thread gets an "invalid event" in the BlockingQueue, it will be terminated.

In the **MultiThreadConsumer class**, I can test the amount of throughput by changing the value of x, the number of threads added in phase two. This adding threads operation is triggered by using CountdownLatch, latch(1), following the requirement of assignment 1.

## Client part 2

Following my client part 1 design, a **CSVWriter** class is added. The required record per request is generated, encapsulated in a new **RequestMetricClass**, and stored in a ConcurrentLinkedQueue. After all posts are sent and all request records are stored in ConcurrentLinkedQueue, the CSVWriter class is called to generate all records in a CSV file. The **MetricCalculator** class handles the calculation of different properties of latency.



# Output

## Client part 1

## SingleThreadConsumer

```
-----  
Single thread consumer test starts  
-----  
Number of successful requests: 10000  
Number of fail requests: 0  
Wall Time: 347975  
Throughout: 34.7975 ms/request  
Throughout: 28 requests/second  
-----  
Single thread consumer test ends  
-----
```

The average response time from a single-thread client is 35 ms/request.

### Multi threads client - part 1

Several options for the number of threads in phase two are tested. The best throughput outcome is 3846 requests/second, with 800 threads in phase two.

According to Little's Law, the predicted throughput is  $200/35 \times 1000 = 5714$  requests/second, which is larger than the test.

Tomcat serve has max 200 threads, but when increasing the client side threads number from 128, 256, 512, to 800, the throughput output of the multi-threads client keeps increasing. Given the Little Law, the minimal value between the number of threads on the server side and the client side, which is 200, should be the bound of the throughput of my program. However, the test of 256, 512, and 800 threads on the client side generated different throughput.

One possible explanation could be more number of client-side threads ensures that the server has a constant queue of requests to process, maximizing its utilization. When the number of client-side threads is too close to the server's capacity, any temporary dip in server-side processing (due to GC pauses, CPU scheduling, etc.) can lead to underutilization. More client threads mean the server is more likely to be kept busy, up to the point where the overhead of managing these additional threads or other resource constraints outweighs the benefits.

```
-----  
Multi threads consumer test start  
-----  
Number of threads in phase 2: 32  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 217596  
Throughout: 921 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
-----  
Multi threads consumer test start  
-----  
Number of threads in phase 2: 64  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 135131  
Throughout: 1481 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
Number of threads in phase 2: 100  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 98952  
Throughout: 2040 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
-----  
Multi threads consumer test start  
-----  
Number of threads in phase 2: 128  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 87555  
Throughout: 2298 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
Number of threads in phase 2: 256  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 67613  
Throughout: 2985 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
Number of threads in phase 2: 512  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 64763  
Throughout: 3125 requests/second  
-----  
Multi threads consumer test end  
-----
```

```
Number of threads in phase 2: 800  
Number of successful requests: 200000  
Number of fail requests: 0  
Wall Time: 52997  
Throughout: 3846 requests/second  
-----  
Multi threads consumer test end  
-----
```

## Multi threads client - part 2

```
Number of threads in phase 2: 128
Number of successful requests: 200000
Number of fail requests: 0
Wall Time: 88089
Mean latency= 37.279186852854565 ms
Median latency= 34.0 ms
P99 latency= 86.0 ms
Throughout: 2272 requests/second
Min latency= 16
Max latency= 10059
-----
Multi threads consumer test end
-----
```

```
Number of threads in phase 2: 256
Number of successful requests: 200000
Number of fail requests: 0
Wall Time: 63224
Mean latency= 41.55439178886217 ms
Median latency= 36.0 ms
P99 latency= 114.0 ms
Throughout: 3174 requests/second
Min latency= 16
Max latency= 10151
-----
Multi threads consumer test end
-----
```

```
Number of threads in phase 2: 512
Number of successful requests: 200000
Number of fail requests: 0
Wall Time: 60123
Mean latency= 63.24680140738075 ms
Median latency= 51.0 ms
P99 latency= 160.0 ms
Throughout: 3333 requests/second
Min latency= 14
Max latency= 10248
-----
Multi threads consumer test end
-----
```

```

Number of threads in phase 2: 800
Number of successful requests: 200000
Number of fail requests: 0
Wall Time: 55162
Mean latency= 103.64576066206234 ms
Median latency= 72.0 ms
P99 latency= 586.0 ms
Throughout: 3636 requests/second
Min latency= 14
Max latency= 20474
-----
Multi threads consumer test end
-----

```

Number of threads	Throughput - client 1	Throughput - client 2
125	2298	2272
256	2985	3174
512	3125	3333
800	3846	3636

The difference between the number of threads in client 1 and client 2 is within 5%.

The plot of throughput over the number of threads in phase 2 (client 1)

Number of threads	Throughput - client 1
32	921
64	1481
100	2040
125	2298
256	2985
512	3125
800	3846



# Throughput

